

Assignment # 2

Modeling Complex Systems (CS/CSYS 302)

Lily Shapiro, Mahalia Clark, & Thomas O’Leary

Part 1: Cellular Automata (CA)

Research Question

How do tree density and fire spread intensity (spread radius) affect forest fire dynamics?

System

To answer our question, we implemented a stochastic cellular automaton in python to model a forest and simulate forest fire dynamics. See “forest_fire_ca.py” for the model code. The CA consists of a 100 x 100 grid with fixed boundaries and three possible states: empty cells (0), living trees (1), and burning trees (2). We initialize the model with empty cells, then randomly assign some of them to be living trees with probability d — the tree density. We then initialize the four central cells as burning trees. We use synchronous updates to update the system according to the following rules:

- For each cell, we count up the number of burning trees in a Moore neighborhood of radius r .
 - We use that count, N , to calculate the probability, p , that a living tree at the center of the neighborhood will catch fire according to: $p = 1 - (1 - q)^N$, where $q = 0.5$ is the probability that each burning tree ignites a central living tree.
 - If a cell is a living tree, it will become a burning tree with this probability, p , depending on the number of burning trees in its Moore neighborhood. Otherwise it will remain living.
 - If a cell is a burning cell, it will become an empty cell in the next timestep. Empty cells will remain empty.
- Methods

To answer our research question, we ran the model for 5 trials of 100 timesteps each for 9 different tree densities, d , and two different Moore neighborhood radii, r , using the following parameter values:

- $r \in \{1, 3\}$
- $d \in \{0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9\}$

Our goal was to compare the dynamics of our CA with those of the deterministic forest fire CA in the textbook, which used a Moore radius of 1 and had living trees catch fire if there was at least one burning tree in its neighborhood. When that forest fire model was tested on a range of initial tree densities, its behavior changed around a threshold of $d = 0.38$. Below this value, the fire tended to go out, whereas, above this value, the fire tended to spread indefinitely. We compare the dynamics of our model to those found in the book to determine the effect of adding a stochastic update rule and increasing the radius of the neighborhood.

While we initialized the forest’s trees randomly according to a variable density, we always initialized our forest fire by setting the four central cells to be burning trees. This acted as a control to make the model runs more comparable by making it easy to see if the single fire went out or spread to the edges of the grid.

We recorded the cumulative areas in each state at each time point and then we averaged those values across the five stochastic trials for each initial density and radii tested. This allowed us to visualize a summary plot of each of the parameter combinations tested. Since we did not treat burnt trees as a separate state in our model, we back-calculated the number of burnt trees by taking the difference between the number of current empty cells and the number of initial empty cells. We did this to quantify and visualize the fire spread.

Assumptions

- This model assumes that fires burn out in one time step. It is possible that it would be more realistic to allow the fire to continue for some (variable/stochastic) number of time steps, therefore increasing the chance that it would burn to other cells.
- Specifically, in the case of radius 3, the CA assumes that the fire may “jump” over healthy trees. This may or may not be rational. I guess sparks could fly lighting off further patches. But, a real fire is unlikely to behave in the sort of jumping and back-burning way that we are allowing here.
- We used synchronous updates because we believed that it was a decent approximation of how a fire spreads in real life (i.e. very quickly and in all directions).
- We assume there is no prevailing wind to drive the fire in a certain direction. Rather, the fire can spread in all directions with equal probability. We also assume that all our living trees are equally flammable, whereas in reality certain patches of forest or individual trees could be more or less dry or surrounded by more or less flammable undergrowth or deadwood.

Results

For a neighborhood of radius 1, the fire quickly burns out at densities of $d < 0.6$ and barely any of the forest is burned. For $d = 0.6$, the fire began to spread and burn a wider area (Figure 2, *left*) but burned area began to plateau towards the end of the 100 time-steps (Figure 1, *left*). For $d > 0.6$ the fire reliably burned outwards in a ring until it consumed most of the forest (Figure 2, *right*). The area of living trees declined over time to near zero while the area of burned trees rose to replace them (Figure 1, *right*). This tells us that the threshold value of d is somewhere around 0.6 for our stochastic model with radius 1. Investigating individual runs with $d = 0.6$ showed that some fires would go out quickly while others would burn to the final timestep. As one could guess intuitively it takes a much denser forest to reach percolation with our stochastic update than with the book’s deterministic update, which only requires one burning fire within the neighborhood. However, once the radius is increased to $r = 3$ (i.e. the fire-spread intensity) the forest fire will percolate at much lower densities ($d \geq 0.2$).

For a neighborhood of radius 3, the fire burns indefinitely in an outwardly spreading ring for any value of $d \geq 0.2$, burning down most of the forest out to the edges of the grid (Figure 3, *right* — the forest plot looks similar to that in Figure 2, *right*). However, for $d = 0.1$, the fire burns only briefly before going out. It does not go out right away, so the threshold value maybe around 0.1 or between 0.1 and 0.2. The threshold values for our model with radius 1 and 3 could be identified more precisely with more runs at a range of densities around the roughly identified threshold regions. The value of q also has a strong impact on the model dynamics, so exploring different values of q could be another avenue for further investigation.

Radius One

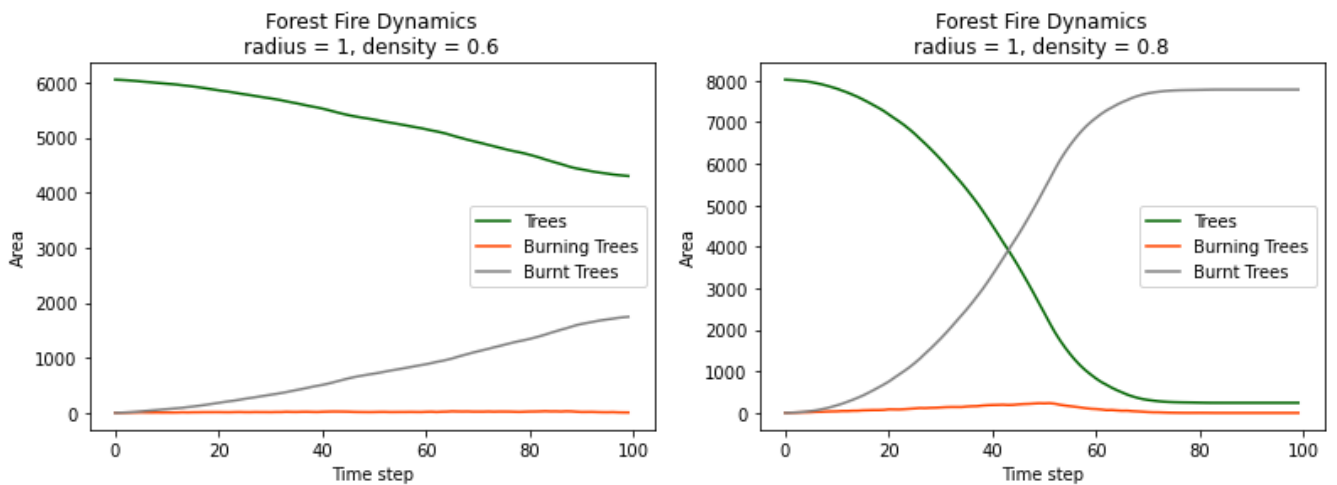


Figure 1: Summary plot of Forest Fire CA overtime with an individual probability of fire of $q = 0.5$ with a neighborhood of radius 1 with differing initial forest densities ($d = 0.6$ left; $d = 0.7$ right). Lines are cumulative areas for the three states: trees, burning trees, and burnt trees (mean of five stochastic trials).




Figure 2: Snapshot of the Forest Fire CA with a radius of one during a run at $t = 49$ with initial densities of $d = 0.6$ *left* and $d = 0.8$ *right*.

Radius Three




Figure 3: Summary plot of Forest Fire CA overtime with an individual probability of fire of $q = 0.5$ with a neighborhood of radius 3 with differing initial forest densities ($d = 0.1$ *left*; $d = 0.2$ *right*). Lines are cumulative areas for the three states: trees, burning trees, and burnt trees (mean of five stochastic trials).

Part 2: Diffusion-Limited Aggregation (DLA)

We implemented a diffusion-limited aggregation (DLA) model in python to simulate the formation of a snowflake from randomly diffusing water vapor particles as they collide with a frozen ice crystal ‘seed’ and freeze onto it. See `Snowflake_DLA.py` for the model code. In our model we have an adjustable n by n grid of cells (where n is odd) with periodic boundaries, where cells can take one of three states: empty (0), randomly floating particle (1), or frozen seed (2).

We initialize the model with empty cells apart from a frozen seed in the central cell. We then add water vapor particles to a random subset of the outermost border of cells, with an adjustable probability d . We use synchronous updates to update the system according to the following rules:

- If a floating particle is touching a frozen cell, it will freeze in its current position.
- If it doesn’t freeze, it will move up, down, left or right with equal probability.
- Frozen cells will stay frozen.

We then add more random particles to the border again after each update, and observe the system periodically. We did not speed it up with gravity or variable step length, so it runs quite slowly.

To get a nice snowglobe effect with swirling particles that can run quickly, we ran the model for 75 timesteps on a 15×15 grid, adding particles with probability $d = 0.1$. However, with such a small grid, individual cells are relatively large, and it is hard to get a pretty fractal, especially as the snowflake tends to freeze out to the edges, locking up the border cells and preventing new particles from drifting in.

To get a nice snowflake fractal we ran the model for 24,000 timesteps on a 55×55 grid, adding snowflakes with a probability $d = 0.0001$, observing the system every 1000 time steps. The larger grid and smaller cells helped to make a lacier fractal that could grow without bumping into edges and freezing the incoming particles.




Figure 3: One of the world’s prettiest digital snowflakes generated by our diffusion-limited aggregation.

Part 3: Gambler's Ruin

Probability of winning the game

Set up

- Suppose you are flipping a biased coin with probability of heads: p .
- You begin the game with an initial $0 < z < W$ dollars. If it is heads you gain a dollar, if it is tails you lose a dollar. The game will end if you run out of money or hit the target of W dollars.
- S_n is the amount of money you have in round n .
- $T = \min\{n : S_n = 0, S_n = W\}$ So T is the number of rounds it takes for you to reach \$0 or \$ W , which is when the game ends

What is the probability that you will eventually win? Prove your conclusion with rigorous arguments.

To begin, similar to the random walk in class, we must first consider the probability of the first step. If P_z is the probability of having your initial value of z then the probability of being one dollar below or above can be determined by p and $(1 - p)$:

$$P_z = pP_{z+1} + (1 - p)P_{z-1}$$

To solve for the chance that you have gained a dollar $P_{z+1} - P_z$ you may substitute P_z as $pP_z + (1 - p)P_z$.

$$\begin{aligned} pP_z + (1 - p)P_z &= pP_{z+1} + (1 - p)P_{z-1} \\ pP_{z+1} - pP_z &= (1 - p)P_z - (1 - p)P_{z-1} \\ p(P_{z+1} - P_z) &= (1 - p)(P_z - P_{z-1}) \\ P_{z+1} - P_z &= \frac{1 - p}{p}(P_z - P_{z-1}) \end{aligned}$$

We can simplify this for $P_2 - P_1$ using the fact that $P_0 = 0$. If we then simplify for $P_3 - P_2$ and substitute in $(P_2 - P_1)$, we can see a recursive pattern, which can be generalized as:

$$\begin{aligned} P_{z+1} - P_z &= \left(\frac{1 - p}{p}\right)^z P_1 \\ P_{z+1} - P_z &= \sum_{n=1}^z (P_{n+1} - P_n) \\ &= \sum_{n=1}^z \left(\frac{1 - p}{p}\right)^n P_1 \\ P_{z+1} &= P_1 + P_1 \sum_{n=1}^z \left(\frac{1 - p}{p}\right)^n \\ &= P_1 \sum_{n=0}^z \left(\frac{1 - p}{p}\right)^n \end{aligned}$$

Since this is a geometric series, we can use a geometric series equation to simplify it to the following:

$$P_{z+1} = P_1 \frac{1 - \left(\frac{1 - p}{p}\right)^{z+1}}{1 - \frac{1 - p}{p}}$$

We can then consider a value of z just one below the target W so that $z = W - 1$. And using the definition $P_W = 1$ (*i.e.* the probability of winning at W is 1 by definition) we can get the equation:

$$P_W = 1 = P_1 \frac{1 - (\frac{1-p}{p})^W}{1 - (\frac{1-p}{p})}$$

We used this to solve for P_1 , which is as follows:

$$P_1 = \frac{1 - \frac{1-p}{p}}{1 - (\frac{1-p}{p})^W}$$

This expression for P_1 can be substituted into the equation for P_{z+1} and simplified to get the following:

$$P_z = \frac{1 - (\frac{1-p}{p})^z}{1 - (\frac{1-p}{p})^W}$$

Probability of a never-ending game

Prove that $Pr(T = \infty) = 0$

$$\begin{aligned} Pr(T = \infty) &= \frac{\# \text{ of infinite paths}}{\text{all paths}} \\ &= \frac{\# \text{ of paths that never reach 0 or W}}{\text{all paths}} \\ &= \frac{(\text{all paths}) - (\# \text{ winning paths}) - (\# \text{ losing paths})}{\text{all paths}} \end{aligned}$$

Probability of winning is > 0 if i is *not equal to* 0 as we saw in the previous question

numerator $<$ all paths

Denominator = All paths = 2^T

As $T \rightarrow \infty$: denominator = $2^T \rightarrow \infty$, and the whole fraction $\rightarrow 0$

So $Pr(T = \infty) = 0$

This shows that our friend is either very generous or very devious depending on how the coin is biased.